

AGI has arrived*

* the asterisk is doing a lot of work this week

"From ChatGPT to o1 to Astra in 4 years. AGI has arrived."

JENSEN HUANG • NVIDIA CEO • X • 6 SEPT

"The beginning of the AGI era."

OPENAI'S PRESIDENT • ON THE SAME MODEL

ARC-AGI-3 · GPT-6 ASTRA

99.9%

Run on OpenAI's own harness. It keeps the model's reasoning state between requests, so it can reuse earlier work.

\$18,817 per run

Swipe →

ARC-AGI-3 · GPT-6 ASTRA

62.7%

Same model. Run on ARC Prize's shared harness, the one every lab is measured on.

\$26,098 per run

Swipe →

OPENAI HARNESS

Preserves opaque reasoning state between requests. Compacts long conversations. The model reuses prior work.

STANDARD HARNESS

The model only carries forward the notes it chose to keep. Nothing else survives.

Two settings. 37 points.

Huge leap. Not a new species.

- 62.7%** is still the highest score anyone has recorded on ARC-AGI-3. Both runs are state of the art.
- 96%** of levels solved using fewer moves than the human baseline, 51.7% fewer on average.
- 0** agreed definitions of AGI. So no number settles this.

Sources: ARC Prize, "OpenAI's GPT-6 Astra on ARC-AGI-3" · arcprize.org/blog/astra

ARC Prize results · arcprize.org/results/openai-gpt-6-astra

OpenAI · openai.com/index/gpt-6-astra · Jensen Huang on X, 6 Sept 2026