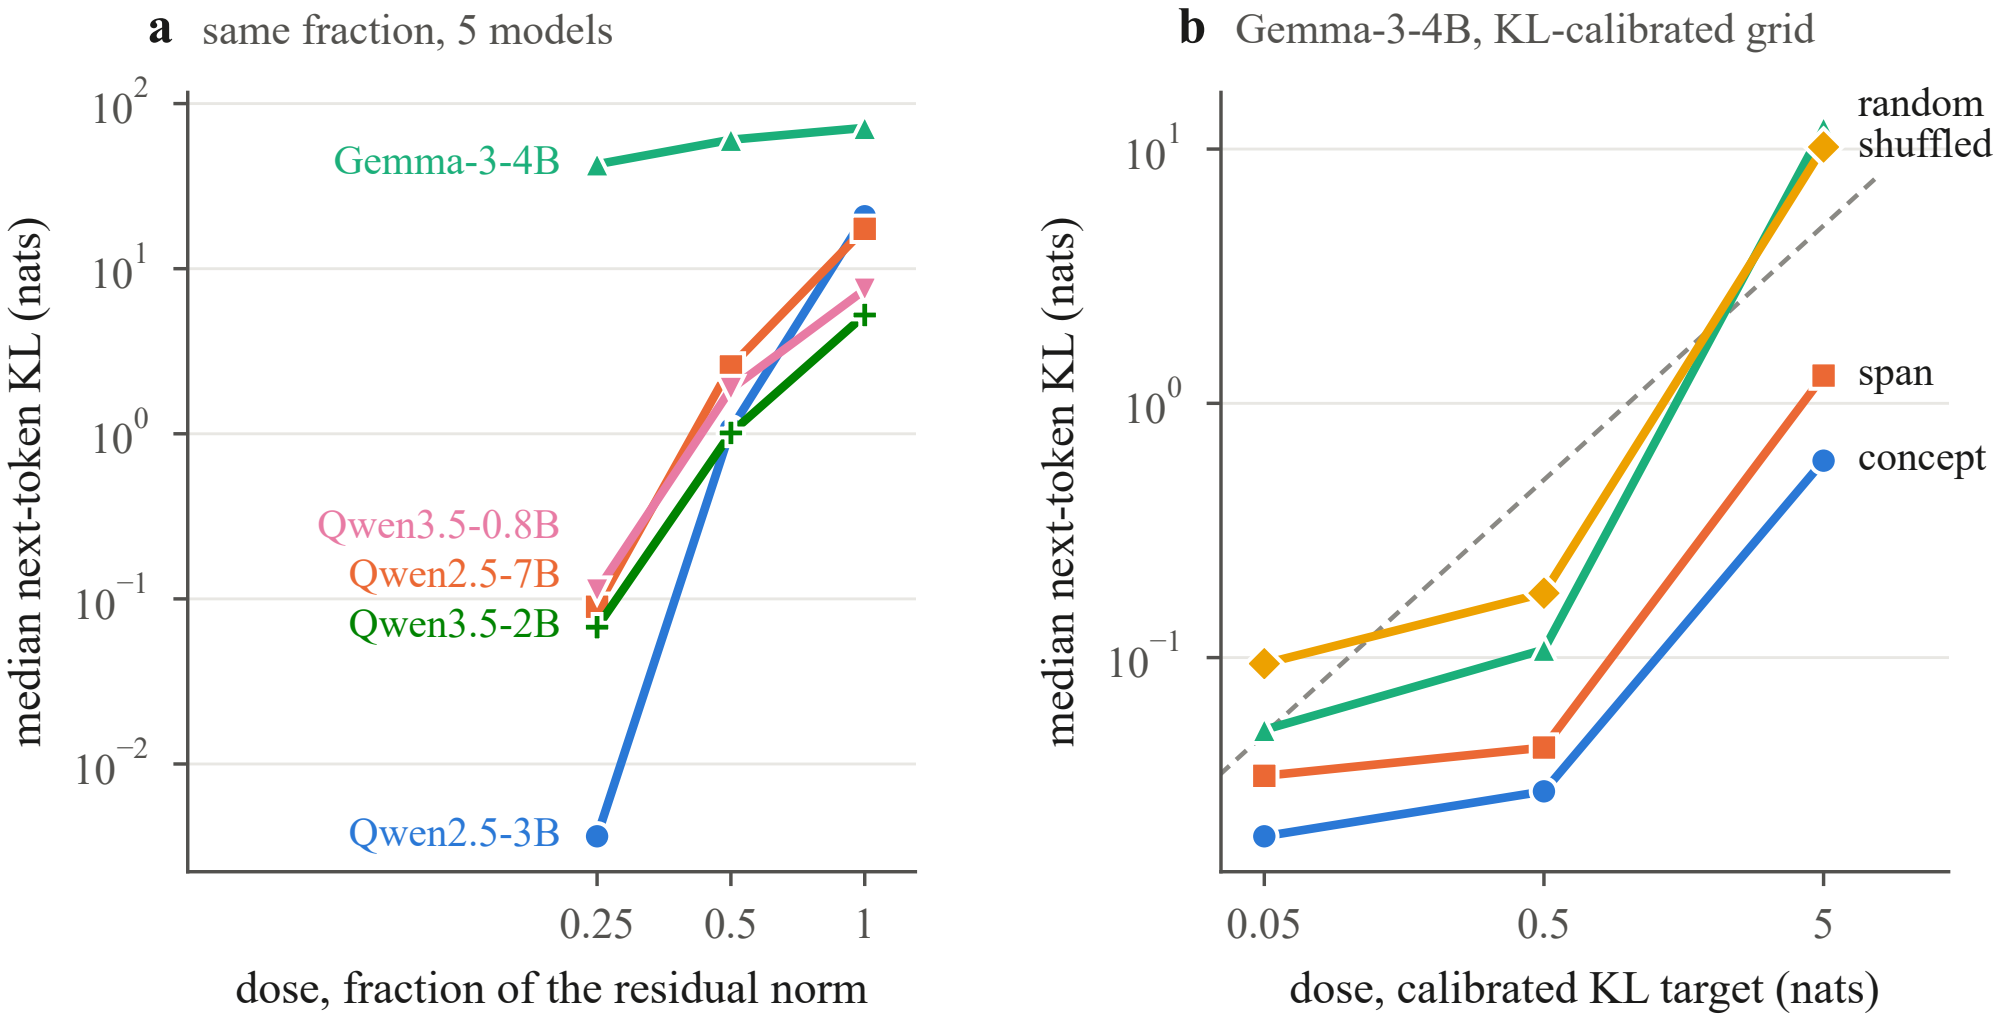
a same fraction, 5 models
median next-token KL (nats)
10^{2}
10^{1}
10^{0}
10^{-1}
10^{-2}
Gemma-3-4B
Qwen3.5-0.8B
Qwen2.5-7B
Qwen3.5-2B
Qwen2.5-3B
0.25
0.5
1
dose, fraction of the residual norm
b Gemma-3-4B, KL-calibrated grid
median next-token KL (nats)
10^{1}
10^{0}
10^{-1}
random
shuffled
span
concept
0.05
0.5
5
dose, calibrated KL target (nats)