


At the same KL, models steer at different rates



the vector's own next-token KL (nats, bin median)

—●— Qwen2.5-3B

—◆— Gemma-3-12B (4-bit)

—■— Qwen2.5-7B

—▼— Gemma-3-27B (4-bit)

—▲— Gemma-3-4B (KL grid)

—★— Gemma-3-27B, released recipe (S-1M)